

Mind the Cap: Output-Budget Regimes Change the Measured Multilingual Reasoning Gap

Anonymous ACL submission

Abstract

Multilingual evaluations often report that models score higher when a problem is first translated into English, but this gap is usually measured at a single output-token cap. We test whether the native-vs-translate gap on MGSM (German, Thai, and Swahili) is a token-budget artifact, using Qwen3-8B and Llama-3.1-8B-Instruct. Budgets are prefixes of one stored 4096-token generation, so token-matched, FLORES-200-length-normalized, and dollar-matched comparisons all derive from a common ledger. Equation (1) shows that the length-normalized contrast is exactly the NATIVE accuracy gain between the matched and the premium-scaled budget, so it must vanish once native accuracy saturates. At the prospectively frozen 1024-token budget, all six Holm-family Qwen tests fail to reject, because native-prefix accuracy has already saturated at moderate levels; Llama likewise fails to reject, but its native accuracy is near floor and parseable answers often never appear. A retrospective sweep locates the budget-binding regime: tighter budgets change the measured gap by up to 38.9 points, and length normalization can reverse which strategy scores higher. COMET translation quality is high but uneven across model-language cells. Above saturation, the residual difference is a strategy-performance gap, not an identified reasoning deficit. The output cap is therefore part of the evaluation estimand: multilingual evaluations should report accuracy across the budget regime, not at a single budget.

1 Introduction

A recurring finding in multilingual LLM evaluation is that models solve reasoning problems better when the problem is first translated into English than when they reason in the original language. This “multilingual reasoning gap” is usually measured at one generation length. Languages differ in how many tokens they need to express paral-

lel content, and strategies differ in when they emit their answer. Under a hard output cap, both features interact with the cap in ways that can resemble a strategy-performance difference.

We ask a narrow, falsifiable question: is the native-vs-translate gap on MGSM a token-budget artifact? We operationalize a budget artifact as the change in that gap between a token-matched cap and a length-normalized, FLORES-200-premium-adjusted cap. If giving NATIVE its premium-scaled budget closes the gap, token matching contributed to the measured difference; if the gap persists, this particular correction does not explain it.

Our contribution is methodological:

1. The measured contrast can depend strongly on the evaluation budget. We pair a prospectively frozen non-rejection at one budget with a retrospective sweep that locates the budget-binding regime.
2. Under tight caps, length normalization can reverse which strategy looks better.
3. Five ledger-computable audits—trace-language ID, COMET translation quality, parser robustness, decoder parity, and normalizer sensitivity—bound the interpretations the stored ledger supports.

Prompt language, reformulation, answer-format compliance, translation quality, and reasoning-trace language are confounded in this design. The object of inference is therefore strategy performance under controlled prefix budgets, not causal reasoning ability.

2 Related work

This study connects multilingual chain-of-thought evaluation on MGSM (Shi et al., 2023) with work on test-time compute, output budgeting, and budget-forcing, including s1 (Muennighoff

et al., 2025). Our length proxy follows the FLORES parallel-text lineage and NLLB’s FLORES-200 benchmark (Goyal et al., 2022; NLLB Team et al., 2022), while the measurement checks use GlotLID (Kargaran et al., 2023) and COMET (Rei et al., 2022). Cross-lingual tokenization disparities make output length a language-specific computational cost (Ahia et al., 2023; Petrov et al., 2023); we study how a prefix budget changes a multilingual strategy contrast after language-length normalization.

Inference-budget work likewise finds that truncating reasoning harms accuracy in model- and modality-dependent ways (Su et al., 2026), while training models to recover answers from truncated traces could shrink the budget-binding regime (Silvestri and Cetin, 2026). Multilingual studies find that English-pivot reasoning can improve accuracy yet drift off target (Luo et al., 2025), and identify input comprehension as a major loss source addressed by an English-intermediary pipeline (Ko et al., 2025). Training-side methods further shrink native-pivot gaps to roughly 2–3.5% under matched supervision (Lasbordes et al., 2026) and produce cross-lingual self-distillation gains that grow with inference budget (Liu et al., 2026). Our contribution is orthogonal and evaluational: these methods should be measured under budget sweeps that reveal whether supervision, strategy, or answer-emission timing drives the reported gap.

3 Design and estimands

Data and strategies. MGSM has 250 items per language in German (de), Thai (th), and Swahili (sw), three languages that span a wide range of FLORES-200 token premiums. We evaluate four strategies, also called arms: NATIVE (reason and answer in language L); TRANSLATE-ACT (translate the problem to English, then solve); PIVOT (English reasoning, answer in L); and CODE-SWITCHED (English scaffold). We draw $k = 8$ samples per item from Qwen3-8B (Qwen Team, 2025) (confirmatory) and Llama-3.1-8B-Instruct (Grattafiori et al., 2024) (secondary, not a replication), for 48,000 stored generations.

Prefix-defined budgets. Each (item, sample) has one stored 4096-token generation. We call this store of generations, together with the scores of all their prefixes, the *ledger*. Evaluating budget B means scoring the length- B prefix. This removes between-budget sampling

noise but limits conclusions to prefix-defined evaluation; independently capped generations may differ. The retrospective regime sweep uses $B \in \{64, 128, 192, 256, 384, 512, 768, 1024\}$; 2048 and 4096 are added only for extended crossover and best-arm checks.

Scoring. We use strict prefix-only exact match on ##### <integer> under intention to treat; truncated, non-integer, and non-compliant answers score 0. Each of the eight samples per item is scored independently at every budget—not by best-of- k , pass@ k , or majority vote—and reported accuracy averages all item-sample cells; the item-clustered bootstrap resamples the 250 MGSM items while retaining all eight samples within each selected item.

Estimand. Let acc_N and acc_T denote NATIVE and TRANSLATE-ACT accuracy, let $\text{gap}(B) = \text{acc}_T(B) - \text{acc}_N(B)$, and let $r_{m,L}$ be the FLORES-200 token premium of language L over English for model m . Then

$$\begin{aligned} \Delta_L(B) &= [\text{acc}_T(B) - \text{acc}_N(B)] \\ &\quad - [\text{acc}_T(B) - \text{acc}_N(\lfloor r_{m,L}B \rfloor)] \\ &= \text{acc}_N(\lfloor r_{m,L}B \rfloor) - \text{acc}_N(B). \end{aligned} \quad (1)$$

TRANSLATE-ACT cancels. What remains is simply how many more NATIVE answers become correct between the matched budget B and the premium-scaled budget $\lfloor r_{m,L}B \rfloor$. We call a budget *binding* when correct answer lines are still appearing inside that interval, and the *budget-binding regime* is the range of B over which they are.

This identity has three direct consequences. First, $\Delta_L(B)$ is a finite, discrete increment of the NATIVE accuracy curve, not a comparator interaction. Second, its peak location follows the NATIVE answer-emission distribution, while its height depends on the premium-scaled interval and the native curve inside it. Third, $\Delta_L(B) \rightarrow 0$ once NATIVE accuracy saturates, so a near-zero value above saturation follows analytically from the estimand.

Primary test. The frozen evaluation budget is $B^* = 1024$, the largest $B \in \{512, 1024\}$ with every $\lfloor r_{m,L}B \rfloor \leq 4096$. The confirmatory family contains six Holm-corrected Qwen tests: H1-existence, directional H1-SESOI ($\Delta_L(B^*) > 5$ points), H2 (a larger effect for Thai than German), and H3 (a dollar-matched strategy reversal) for each of the three languages. Inference uses

an item-clustered paired bootstrap (10,000 resamples), a studentized sup- t maximum statistic, and a pre-specified $1.3\times$ tail-conservatism factor; Appendix D gives details. The exploratory two-sided equivalence statement below is separate. The analysis protocol was internally frozen before confirmatory scoring; this was not a public preregistration.

FLORES-200 premiums are Qwen de 1.56 / th 2.55 / sw 1.94 and Llama de 1.58 / th 2.19 / sw 1.93. For each model and language, $r_{m,L}$ is the total-token ratio to parallel English over all 1,012 FLORES-200 devtest sentence pairs after NFC normalization, computed with that model’s own tokenizer and a paired percentile bootstrap over sentence pairs.

4 Regime-dependent results

4.1 The frozen test yields no confirmatory support

At $B^* = 1024$, $\Delta_L(B^*)$ is near zero for every Qwen language, and all six Holm-family tests fail to reject (formal outcome no_confirmatory_h1_support):

- Qwen $\Delta_L(B^*)$: de 0.00, th 0.15, sw 0.05 points. H1-existence raw $p = 0.060$ at Holm-local $\alpha = 0.0083$; H1-SESOI raw $p = 1.0$.
- Llama $\Delta_L(B^*)$: de/th/sw all 0.00. This procedurally matched secondary analysis is outside the confirmatory family and also rejects nothing.

The two non-rejections have different descriptive causes. Qwen NATIVE accuracy plateaus at 79.0% / 47.1% / 33.7% for de/th/sw; between 1024 and the Thai FLORES-scaled budget of 2611, only 3 of 2000 traces become newly correct. Llama NATIVE instead plateaus near floor at 13.6% / 3.9% / 29.0%, alongside never-emission rates of 80.2% / 93.2% / 46.0%. Qwen therefore reaches score saturation at moderate accuracy, whereas Llama de/th often never produces a parseable native answer. Neither pattern implies that all traces have terminated.

4.2 Tight budgets expose a large, then vanishing, artifact

Table 1 places the largest retrospective $\Delta_L(B)$ inside the budget-binding regime (pointwise item-clustered bootstrap 95% CIs).

The Qwen rows of Table 1 contain three different kinds of peak. German gives the clearest two-arm comparison: at $B = 192$ both arms are well away from floor, and premium-scaling raises NATIVE by 34.2 points.

Thai has the largest peak, and also the largest premium: at $B = 256$, $r = 2.55$ gives a premium-scaled budget of 652, so $\Delta_L(256)$ summarizes the gain across a 396-token interval rather than a contrast at equal output widths.

Swahili is weaker evidence. Its peak occurs where TRANSLATE-ACT scores only 0.60%, so it shows a native-prefix rescue rather than a clean two-arm comparison away from floor. The same caution applies to the German crossover at $B = 128$ (§4.3).

Simultaneous max- $|t|$ 95% bands over the sweep keep every Qwen peak away from zero (de@192 [27.8, 40.6], th@256 [32.3, 45.4], sw@128 [10.7, 19.3]), and peak locations are stable in 89.6%, 100.0%, and 87.9% of bootstrap replicates. At B^* the largest language-specific upper bound is 0.32 points for Qwen and 0.00 for Llama, below the exploratory ± 5 -point equivalence margin.

Normalizer sensitivity. The behavioral trace-length ratio is smaller than the FLORES premium in all six model-language cells, with non-overlapping pointwise intervals (Appendix C). FLORES therefore consistently grants the larger premium.

Substituting the behavioral ratio would still produce a 5-point artifact in four of the six cells. The minimum evaluated premium that produces a 5-point artifact lies below the behavioral ratio for Qwen de (1.089 vs 1.467) and th (1.188 vs 2.038), and for Llama de (1.274 vs 1.432) and sw (1.253 vs 1.848). Qwen Swahili is the exception: its threshold of 1.254 sits above the behavioral ratio of 1.179, so the 5-point Swahili claim depends on a premium above that ratio, including the frozen FLORES token premium. Llama Thai never reaches 5 points, even at $1.5\times$ FLORES.

These thresholds are grid-selected and conditional on the stored traces. The behavioral ratio is not a validated normalizer, because it compares traces that differ in content, correctness, and stopping behavior.

Figure 1 plots the Qwen NATIVE curves and marks both accuracy endpoints of each peak premium interval.

Table 1: Peak length-normalization effects, where $\Delta_L(B) = \text{acc}_N(\lfloor rB \rfloor) - \text{acc}_N(B)$, together with values at the frozen budgets. Intervals are pointwise item-clustered bootstrap 95% CIs.

model/lang	peak $\Delta_L(B)$ (pts)	peak B	NATIVE acc. at B	TRANSLATE-ACT acc. at B	$\Delta_L(512)$	$\Delta_L(1024)$
Qwen de	34.2 [30.2, 38.3]	192	16.10	22.60	2.25	0.00
Qwen th	38.9 [34.7, 43.0]	256	6.20	47.75	8.85	0.15
Qwen sw	15.0 [12.4, 17.7]	128	8.70	0.60	0.25	0.05
Llama de	8.4 [7.0, 9.9]	256	3.85	43.10	0.15	0.00
Llama th	2.3 [1.6, 3.1]	192	1.10	16.90	0.15	0.00
Llama sw	18.2 [15.9, 20.6]	256	8.25	44.60	1.60	0.00




Figure 1: Qwen NATIVE accuracy by budget. Shaded intervals mark each language’s peak normalization interval ($B_{\text{peak}}, \lfloor rB_{\text{peak}} \rfloor$); the vertical gain between the endpoint markers equals the peak $\Delta_L(B)$.

4.3 Tight caps can reverse strategy rankings

Table 2 shows that Qwen crossover strength follows the left tail of answer-emission timing more closely than the medians.

TRANSLATE-ACT’s median emission is earlier than NATIVE’s in German and Thai. The p10 ordering is timing-consistent with the crossover: Swahili has a 74.7-token NATIVE early-tail advantage and the strongest native lead; German differs by only 2 tokens and has a narrow lead; Thai NATIVE is 71 tokens later and never leads. These grid-resolved emission summaries do not isolate translation-segment length or establish mediation.

At $B = 128$, Qwen German NATIVE scores 2.55% versus 1.15% for TRANSLATE-ACT; Swahili NATIVE leads with bootstrap probability 1.00 at 128 and 192 (transition [192,256]), while German leads with probability 0.958 at 128 (transition [128,192]). Thai has no native lead, and no Llama language crosses because native accuracy remains near floor. The Qwen Swahili crossover also comes from a degenerate heavy-tail cell: NATIVE output-length p90 is 4096 and 25.1% never emit a parseable answer. The frozen H3 test did not detect these reversals because it examined only the looser pre-specified budgets.

5 Measurement audits

We report five ledger audits and one diagnostic. Trace-language ID, COMET, parser robustness, and decoder parity appear below; normalizer sensitivity appears in §4.2. The verbosity/failure-tail decomposition is a diagnostic, not a sixth audit.

Trace-language ID (automated GlotLID). Automated labels support the main NATIVE vs TRANSLATE-ACT contrast. Determinate NATIVE traces are classified in L at Qwen de 92.1% / sw 94.1% / th 99.4%, and at approximately 100% for all Llama languages; TRANSLATE-ACT post-delimiter reasoning is 98.4–99.9% English.

Three caveats apply. First, the initial swh-only mapping failed the frozen native:sw agreement criterion ($75\% < 90\%$). We then adopted the Swahili macrolanguage remap (swh+swc, with swc denoting Congo Swahili), which raised Qwen native-Swahili compliance from 85.8% to 94.1% and the validation cell to 90.00% (18/20); because this remap was adopted after inspecting the failure, we report it as a post-hoc analytic decision. Unrelated neighboring Bantu labels remain outside Swahili. Second, an independent blind LLM adjudication of 240 Qwen traces agrees at 96.7% overall and at least 90% in every cell, but it is preliminary and the frozen human validation remains outstand-

Table 2: Qwen answer-emission timing and observed strategy crossovers. E is the token position at which a parseable answer is emitted; p10 is its 10th percentile.

lang	NATIVE median E	TRANSLATE-ACT median E	NATIVE p10 E	TRANSLATE-ACT p10 E	observed crossover
sw	206	262	96	170.7	strongest: NATIVE leads at 64/128/192
de	270	247	160	158	marginal: NATIVE leads at 128
th	377	250	236	165	none

ing. Third, PIVOT and CODE-SWITCHED violate their English instruction in 9 of 12 cells, so they remain outside the main comparison.

COMET translation quality. Reference-based COMET is generally high but nonuniform (Table 3): German is strongest for both models (0.877 Qwen, 0.872 Llama), Qwen Swahili is lowest at 0.749, and Llama Thai has a weak lower tail (p10 0.325). At the six prespecified peak budgets, an exploratory per-cell analysis finds weak and inconsistent COMET-correctness-gain associations (Spearman $\rho = -0.143$ to 0.280, five positive, only one with $|\rho| \geq 0.20$; Table 4), so the tight-budget TRANSLATE-ACT advantage is not well explained by translation quality. Translation quality remains part of the strategy bundle; these scores are descriptive and never condition accuracy.

Parser robustness. In NATIVE peak cells, prefix-only rescued-correct traces are at most 0.35% and value-unstable traces at most 0.30%. Within the $(B, \lfloor rB \rfloor]$ intervals, 96.8–100% of native gains are genuinely terminated, and requiring a terminated answer line changes every peak by at most 0.2 points. The tight-budget effect is therefore late answer emission rather than a prefix-parser artifact.

Decoder parity. A stratified 2,520-prefix audit finds 37.9% raw exact-string agreement between local and vLLM decoding but 100% agreement after production special-token normalization, including 100% parsed-answer and correctness agreement. All raw divergences are cosmetic special-token markup.

Verbosity/failure-tail diagnostic. Qwen Swahili NATIVE has a heavy tail: 10.6% hit the 4096 cap and 25.1% never emit a parseable answer. This mixes truncation, non-integer or multiple answers, and format noncompliance, cautioning against interpreting native accuracy as pure reasoning ability.

6 Scope and implications

In these stored-prefix MGSM evaluations, multilingual exact-match comparisons are regime-dependent: length normalization has large descriptive effects while answer emission binds, and negligible effects after score saturation. The prospectively frozen non-rejection at $B^* = 1024$ is a negative boundary condition; the retrospective sweep locates where the sensitivity actually occurs.

At saturated budgets the residual native-vs-translate difference is large (Qwen Thai about +41 points, Llama Thai about +69), but it is a strategy-performance gap. Prompt language, reformulation, format compliance, translation quality, and reasoning language remain confounded.

Calling this residual “not an identified reasoning deficit” states what our design can identify, not that execution effects are absent: complementary diagnostics find genuine target-language reasoning-execution failures even with English inputs (Zhang et al., 2026).

Limitations. (i) Scores come from prefixes of one long decode, not independently hard-capped decodes; decoder parity establishes scoring parity, not trajectory parity, and hard caps could elicit different trajectories, so the prefix-based magnitudes are an upper-relevance bound pending independently capped replication. The tight-budget result is retrospective and exploratory. (ii) vLLM repeat generation was only 46% bitwise deterministic, preserving the stored-ledger estimand but weakening reproducibility and shared-seed interpretation. (iii) Scope is MGSM, three languages, and two 8B models. (iv) The frozen human GlotLID validation remains outstanding, despite the preliminary blind LLM agreement. (v) The frozen same-content trace-premium validation was not performed; the behavioral ratio in Appendix C does not substitute for it. (vi) Calibration is only approximately nominal (corrected type-I 0.00917 vs 0.00833 target); the $1.3\times$ factor is a family-wide safeguard, not a verified family-wise calibration.

Future work should test prompt interventions

that elicit earlier answer emission—which would shrink NATIVE’s budget-binding regime—and extend the protocol beyond MGSM numeric exact match to longer-form and non-numeric tasks.

Data and code availability. The full ledger of stored generations and all analysis code will be released.

Takeaway. When comparing language strategies under a token cap, treat the cap as an independent variable rather than a background constant. Report accuracy across a budget sweep with a length-normalized comparison, mark the region in which answer emission binds, and do not let a single budget carry the claim.

References

Orevaoghene Ahia, Sachin Kumar, Hila Gonen, Jungo Kasai, David R. Mortensen, Noah A. Smith, and Yulia Tsvetkov. 2023. Do all languages cost the same? tokenization in the era of commercial language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 9904–9923.

Naman Goyal, Cynthia Gao, Vishrav Chaudhary, Peng-Jen Chen, Guillaume Wenzek, Da Ju, Sanjana Krishnan, Marc’Aurelio Ranzato, Francisco Guzmán, and Angela Fan. 2022. The flores-101 evaluation benchmark for low-resource and multilingual machine translation. *Transactions of the Association for Computational Linguistics (TACL)*, 10:522–538.

Aaron Grattafiori et al. 2024. *The llama 3 herd of models*. *arXiv preprint arXiv:2407.21783*.

Amir Hossein Kargaran, Ayyoob Imani, François Yvon, and Hinrich Schütze. 2023. GlotLID: Language identification for low-resource languages. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 6155–6218.

Hyunwoo Ko, Guijin Son, and Dasol Choi. 2025. *Understand, solve and translate: Bridging the multilingual mathematical reasoning gap*. *arXiv preprint arXiv:2501.02448*.

Maxence Lasbordes, Amélie Chatelain, and Djamel Seddah. 2026. *Rethinking the multilingual reasoning gap with layer swap*. *arXiv preprint arXiv:2605.26735*.

Yihong Liu, Raoyuan Zhao, Michael A. Hedderich, and Hinrich Schütze. 2026. *Crosslingual on-policy self-distillation for multilingual reasoning*. *arXiv preprint arXiv:2605.09548*.

Wenyang Luo, Wayne Xin Zhao, Jing Sha, Shijin Wang, and Ji-Rong Wen. 2025. *MMATH: A multilingual benchmark for mathematical reasoning*. *arXiv preprint arXiv:2505.19126*.

Niklas Muennighoff, Zitong Yang, Weijia Shi, Xiang Lisa Li, Li Fei-Fei, Hannaneh Hajishirzi, Luke Zettlemoyer, Percy Liang, Emmanuel Candès, and Tatsunori Hashimoto. 2025. *s1: Simple test-time scaling*. *arXiv preprint arXiv:2501.19393*.

NLLB Team, Marta R. Costa-jussà, James Cross, Onur Çelebi, Maha Elbayad, Kenneth Heafield, Kevin Heffernan, Elahe Kalbassi, Janice Lam, Daniel Licht, Jean Maillard, et al. 2022. *No language left behind: Scaling human-centered machine translation*. *arXiv preprint arXiv:2207.04672*. Introduces the FLORES-200 benchmark.

Aleksandar Petrov, Emanuele La Malfa, Philip Torr, and Adel Bibi. 2023. Language model tokenizers introduce unfairness between languages. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 36.

Qwen Team. 2025. *Qwen3 technical report*. *arXiv preprint arXiv:2505.09388*.

Ricardo Rei, José G.C. de Souza, Duarte Alves, Chrysoula Zerva, Ana C. Farinha, Taisiya Glushkova, Alon Lavie, Luisa Coheur, and André F.T. Martins. 2022. COMET-22: Unbabel-IST 2022 submission for the metrics shared task. In *Proceedings of the Seventh Conference on Machine Translation (WMT)*, pages 578–585.

Freda Shi, Mirac Suzgun, Markus Freitag, Xuezhi Wang, Suraj Srivats, Soroush Vosoughi, Hyung Won Chung, Yi Tay, Sebastian Ruder, Denny Zhou, Dipanjan Das, and Jason Wei. 2023. *Language models are multilingual chain-of-thought reasoners*. In *The Eleventh International Conference on Learning Representations (ICLR)*.

Gianluigi Silvestri and Edoardo Cetin. 2026. *Learning from partial chain-of-thought via truncated-reasoning self-distillation*. *arXiv preprint arXiv:2603.13274*.

Ian Su, Gaurav Purushothaman, Jey Narayan, Ruhika Goel, Kevin Zhu, Sunishchal Dev, Yash More, and Maheep Chaudhary. 2026. *Broken chains: The cost of incomplete reasoning in LLMs*. *arXiv preprint arXiv:2602.14444*.

Jiaqiao Zhang, Zhoujun Li, Raoyuan Zhao, Jian Lan, Thomas Seidl, Michael A. Hedderich, Hinrich Schütze, and Yihong Liu. 2026. *Beyond input understanding: Diagnosing multilingual mathematical reasoning with directed acyclic trace graphs*. *arXiv preprint arXiv:2605.27715*.

A Accuracy-curve summary

The complete curves use $B \in \{64, 128, 192, 256, 384, 512, 768, 1024\}$. Qwen NATIVE accuracy rises from 0.0/0.0/0.2% at $B = 64$ to 79.0/47.1/33.7% at $B = 1024$ for de/th/sw; Llama NATIVE rises from 0.0/0.0/0.0%

to 13.6/3.9/29.0%. The Qwen curves climb sharply through the peak intervals in Figure 1 and then flatten, whereas the Llama de/th curves remain near floor. The supplementary material reports every model–language–arm value at all eight budgets.

B Best-English-arm

Reselecting the empirically best English arm inside each bootstrap replicate reproduces the pre-selected TRANSLATE-ACT gap closely (Qwen Thai about 41 points, Llama Thai about 69; uplift over TRANSLATE-ACT is near zero in most cells).

C Trace-length ratio (not a normalizer)

The ratio of median NATIVE output tokens to median TRANSLATE-ACT post-delimiter English-reasoning tokens is Qwen de 1.47 / th 2.04 / sw 1.18. Subtracting the FLORES premium gives a negative difference in all six cells, with non-overlapping pointwise intervals: Qwen de -0.092 $[-0.144, -0.028]$, th -0.513 $[-0.588, -0.440]$, sw -0.757 $[-0.824, -0.685]$; Llama de -0.149 $[-0.193, -0.099]$, th -0.427 $[-0.488, -0.365]$, sw -0.083 $[-0.129, -0.021]$. This ratio compares behaviorally different traces and cannot validate or replace FLORES.

D Statistical machinery and the six-test family

The Qwen confirmatory family contains exactly six raw p-values, corrected by Holm at family level $\alpha = 0.05$: H1-existence tests whether any language has $\Delta_L > 0$; H1-SESOI separately tests whether any has $\Delta_L > 5$ points; H2 tests the directional ordered contrast $\Delta_{th} - \Delta_{de} > 0$; and H3-de, H3-th, and H3-sw each test for a dollar-matched strategy reversal, requiring a significantly positive NATIVE-minus-TRANSLATE-ACT contrast at one common-support budget and a significantly negative contrast at another. Each H3 p-value is the maximum of the multiplicity-controlled positive- and negative-side p-values, an intersection-union test.

Inference resamples the 250 parallel MGSM items as clusters, retaining all languages, arms, and eight samples within each selected item, for 10,000 paired bootstrap resamples. Studentized sup- t maxima provide simultaneous bounds over languages for H1 and over budgets for H3; H2

uses the corresponding one-sided studentized bootstrap contrast. Because calibration with 250 discrete item clusters found mild anti-conservatism in the extreme $\alpha/6$ tail, every raw confirmatory p-value is multiplied by the pre-specified 1.3 tail-conservatism factor (capped at 1), equivalently testing at $\alpha/1.3$, before Holm correction.

E TRANSLATE-ACT translation quality (COMET, descriptive)

Reference-based Unbabel/wmt22-comet-da; one stored full trace per item (sample 0); delimiter-missing traces excluded from this denominator only; pointwise 95% item-bootstrap CIs (10,000 resamples).

These scores are descriptive only and never condition, gate, or reweight accuracy.

Table 3: Reference-based COMET translation quality for TRANSLATE-ACT. Intervals are pointwise 95% item-bootstrap CIs.

Model	Lang	n	Missing delimiter	COMET mean	median	p10	p90	95% CI
Qwen	de	250	0.0%	0.877	0.878	0.834	0.921	[0.872, 0.881]
Qwen	th	249	0.4%	0.858	0.866	0.810	0.907	[0.851, 0.865]
Qwen	sw	250	0.0%	0.749	0.772	0.580	0.867	[0.734, 0.763]
Llama	de	249	0.4%	0.872	0.874	0.826	0.915	[0.868, 0.877]
Llama	th	244	2.4%	0.783	0.850	0.325	0.895	[0.759, 0.805]
Llama	sw	240	4.0%	0.798	0.825	0.703	0.888	[0.783, 0.811]

Table 4: Exploratory COMET association with correctness gain at each prespecified peak token budget, using sample 0. Intervals are pointwise bootstrap 95% CIs.

Model	Lang	Peak B	n	Spearman ρ [95% CI]
Qwen	de	192	250	0.015 [−0.112, 0.139]
Qwen	th	256	249	0.146 [0.019, 0.269]
Qwen	sw	128	250	−0.143 [−0.269, −0.004]
Llama	de	256	249	0.280 [0.155, 0.399]
Llama	th	192	244	0.009 [−0.129, 0.146]
Llama	sw	256	240	0.074 [−0.051, 0.203]